CORPORATE POLICY DOCUMENT

Artificial Intelligence Acceptable Use and Governance Policy

Governing the Use of Large Language Models in Regulated Research, Clinical, and Regulatory Affairs Workflows

Policy Number	POL-AI-001
Version	1.0 — Effective upon Board Approval
Effective Date	[DATE]
Next Review Date	[DATE + 12 months]
Policy Owner	Chief Compliance Officer
Co-Owners	Chief Scientific Officer | General Counsel | VP Regulatory Affairs
Applies To	All Employees, Contractors, Consultants, and Third Parties with access to Company AI Systems
Regulatory Scope	21 CFR Parts 11, 312, 314, 601 | ICH E6(R3) | HIPAA 45 CFR Parts 160 & 164 | EU AI Act (2024/1689) | EU GDPR
Classification	CONFIDENTIAL — Internal Use

LEGAL NOTICE: This Policy constitutes binding corporate governance. Violation may result in disciplinary action up to and including termination of employment or contract, referral to regulatory authorities, and civil or criminal liability. Nothing in this Policy constitutes legal advice. Specific regulatory questions must be directed to the Legal Department or outside counsel.

Approvals and Revision History

Approval Signatures

Role	Signature	Date
Chief Compliance Officer	_____	[DATE]
Chief Scientific Officer	_____	[DATE]
General Counsel	_____	[DATE]
VP Regulatory Affairs	_____	[DATE]
Chief Information Officer	_____	[DATE]
Board Audit Committee Chair	_____	[DATE]

Revision History

Version	Date	Author(s)	Summary of Changes
1.0	[DATE]	Policy Working Group	Initial Policy — Approved by Board
			(Reserved for future revisions)

Section 1 — Purpose and Scope

1.1 Purpose

This Artificial Intelligence Acceptable Use and Governance Policy ("Policy") establishes the standards, obligations, and controls governing the access to, deployment of, and reliance upon artificial intelligence systems — including large language models ("LLMs"), generative AI platforms, and AI-assisted analytics tools — by [Company Name] ("Company" or "we") and all personnel operating within its organizational boundaries.

The Company operates as a regulated entity subject to the oversight of the U.S. Food and Drug Administration ("FDA") under the Federal Food, Drug, and Cosmetic Act ("FD&C Act") and its implementing regulations, including Good Laboratory Practice (GLP) regulations at 21 CFR Part 58, Good Clinical Practice (GCP) standards embodied in 21 CFR Parts 312 and 314, and the electronic records and signatures requirements of 21 CFR Part 11. The introduction of AI systems into regulated workflows creates obligations that are not fully addressed by vendor terms of service and that require the Company to establish internal governance commensurate with the compliance demands of a GxP-regulated entity.

This Policy is further motivated by the FDA's October 2023 discussion paper "Artificial Intelligence in Drug Development: Issues and Considerations" (Docket FDA-2023-N-3379), the FDA's January 2025 AI Action Plan, and the European Union's Artificial Intelligence Act (Regulation (EU) 2024/1689, effective August 2024), each of which signals a regulatory environment in which sponsors and manufacturers are expected to exercise demonstrable oversight and control over AI tools used in regulated activities.

1.2 Scope

This Policy applies to:

- All full-time and part-time employees of the Company regardless of location;
- All contractors, consultants, temporary workers, and agency staff engaged by the Company who access Company systems or Company data;
- All third-party organizations acting as Business Associates under HIPAA or as service providers with access to Company Confidential Information or Company AI Systems;
- All AI Systems, whether operated by the Company internally, accessed through application programming interfaces (APIs), or provided by third-party vendors as Software-as-a-Service (SaaS), that are used in connection with Company activities.

This Policy applies across all functional departments, with additional departmental-specific requirements set forth in Section 5. It covers AI use in, without limitation: scientific research and discovery; preclinical and clinical development; regulatory strategy and submission; pharmacovigilance and safety; intellectual property analysis; legal and compliance functions; commercial and medical affairs; and administrative and operational functions.

Out of Scope

This Policy does not govern AI systems embedded within validated, commercially licensed software platforms that are operated solely through their standard user interfaces with no custom prompting, no export of underlying model outputs to regulated documents, and no integration with Company data beyond standard system inputs (e.g., an AI-powered spell-checker embedded in a word processing platform). The AI Governance Committee may issue guidance clarifying the scope boundary for specific embedded AI features.

Section 2 — Definitions

The following terms have the meanings set forth below when used in this Policy. Terms not defined herein have the meanings ascribed to them in the applicable federal regulations or ICH guidelines.

Term	Definition
AI System	Any software that uses machine learning, deep learning, natural language processing, or other artificial intelligence techniques to perform tasks that would ordinarily require human cognition, including but not limited to generating text, summarizing documents, answering questions, extracting information, classifying data, or producing code. For purposes of this Policy, "AI System" encompasses large language models (LLMs) such as Claude (Anthropic) and GPT-4 (OpenAI), AI-assisted literature review tools, AI-powered coding assistants, and AI-enabled data analytics platforms.
Validated Use	An application of an AI System that has been: (i) documented in a User Requirements Specification (URS); (ii) assessed for risk in accordance with the Company's risk management SOP; (iii) tested and qualified through an Installation Qualification (IQ), Operational Qualification (OQ), and, where required by the risk assessment, Performance Qualification (PQ); and (iv) approved for production use by the AI Governance Committee and, for GxP-applicable uses, by the Quality Assurance function. A Validated Use is distinguished from an exploratory or non-regulated use by the formality of its qualification documentation and the regulatory context in which its outputs are employed.
GxP Context	Any activity governed by Good Practice ("GxP") regulations promulgated by FDA or international regulatory authorities, including Good Laboratory Practice (21 CFR Part 58), Good Clinical Practice (ICH E6(R3); 21 CFR Parts 312, 314), Good Manufacturing Practice (21 CFR Parts 210-211; ICH Q10), Good Pharmacovigilance Practice (ICH E2E), and Good Regulatory Practice. AI outputs that are incorporated into GxP-applicable documents — including laboratory notebooks, clinical trial protocols, investigator brochures, clinical study reports, INDs, NDAs, BLAs, DSURs, and deviation records — are generated in a GxP Context and are subject to all applicable requirements of this Policy.
Protected Health Information (PHI)	Individually identifiable health information, including electronic PHI (ePHI), as defined at 45 CFR § 160.103 under the Health Insurance Portability and Accountability Act (HIPAA) Privacy Rule. PHI includes any information relating to a subject's past, present, or future physical or mental health condition, the provision of health care to the subject, or the past, present, or future payment for health care, that identifies or could reasonably be used to identify the subject. In the clinical trial context, PHI includes, without limitation, subject medical records, adverse event narratives containing identifiable subject information, imaging data, genomic data linked to a subject identifier, and informed consent forms containing subject signatures.
Confidential Company Information (CCI)	Non-public information owned by or in the possession of the Company that has economic value from its secrecy, including trade secrets, proprietary compound structures and synthesis data, unpublished clinical trial data, regulatory strategy documents, financial projections, business development targets, attorney-client privileged communications, and work product. CCI is distinct from PHI but may include PHI in certain contexts (e.g., a clinical study report containing both proprietary efficacy data and subject-level safety narratives).
Human Oversight	The review, evaluation, and approval of AI-generated output by a qualified human being who: (i) possesses the professional knowledge, credentials, and experience necessary to assess the accuracy, completeness, and appropriateness of the output in its intended regulatory or scientific context; (ii) exercises substantive judgment rather than perfunctory sign-off; and (iii) takes professional and personal responsibility for the accuracy of any AI-generated content incorporated into a document under their name or signature.
AI Governance Committee (AGC)	The standing cross-functional committee established under Section 14 of this Policy, composed of representatives from Quality Assurance, Regulatory Affairs, Legal, IT/Informatics, and senior scientific leadership, responsible for oversight of the Company's AI program, approval of new AI use cases, management of the AI System Registry, and disposition of Policy exceptions.
Audit Trail	A secure, computer-generated, time-stamped electronic record that allows

	reconstruction of the course of events relating to the creation, modification, deletion, or transmission of an electronic record, as required by 21 CFR § 11.10(e). In the context of AI use, the Audit Trail captures the provenance of AI-generated content incorporated into regulated records, including the AI System used, model version, timestamp of generation, identity of the human operator, and identity and timestamp of the reviewing authority.
Model Version	A specific, identified release of an AI System's underlying model weights and architecture that determines the system's behavior and output characteristics. Model Versions are relevant to validation because different Model Versions of the same AI System may produce materially different outputs for identical inputs; a change in Model Version constitutes a change in the computerized system requiring assessment under the Company's change control SOP.
Prompt	A natural language instruction, question, or set of instructions submitted to an AI System by a user that elicits a response from the model. In regulated contexts, Prompts that generate content incorporated into GxP-applicable documents must be documented as part of the Audit Trail.

Section 3 — Policy Statement

The Company recognizes that large language models and AI systems offer material opportunities to accelerate scientific discovery, improve documentation quality, and enhance operational efficiency across the drug development lifecycle. The Company is equally committed to the principle that the integrity of its regulated activities — and the safety of patients who depend on its products — cannot be compromised by the uncritical adoption of AI tools whose outputs require expert validation.

Accordingly, the Company adopts the following foundational principles governing all AI use:

1. Human Primacy. AI Systems are tools that augment human expertise; they do not replace the professional judgment of qualified scientists, clinicians, regulatory professionals, attorneys, or compliance specialists. No AI-generated output shall be incorporated into a regulated document, a legal instrument, a patient-facing communication, or a business decision of material consequence without substantive review and approval by a qualified human being.
2. Regulatory Defensibility. All AI use in regulated activities must be designed and documented such that the Company could demonstrate, in the context of an FDA inspection, EMA audit, or legal proceeding, that: (a) the AI System used was appropriate for the intended purpose; (b) its outputs were subject to qualified human oversight; (c) an adequate audit trail exists; and (d) the Company bears full responsibility for the accuracy of the regulated output.
3. Data Protection. Company data — including PHI, CCI, and proprietary scientific data — shall not be transmitted to external AI systems in a manner that creates regulatory, legal, or competitive risk. The protections applicable to a piece of information in any other context apply equally when that information is considered for use as an AI input.
4. Proportionate Controls. Governance requirements are calibrated to the risk profile of the use case. AI applications in non-regulated, non-PHI contexts operate under a lighter framework than AI applications that generate content destined for FDA submissions or that process patient data. The risk-tiering framework is set forth in Appendix A.
5. Continuous Oversight. The regulatory and technical landscape governing AI in pharmaceutical development is evolving. This Policy and its associated procedures shall be reviewed no less frequently than annually and updated to reflect material developments in FDA guidance, ICH standards, EU AI Act implementing acts, and industry best practice.

Section 4 — AI System Registry and Approval Requirements

4.1 Registry Requirement

The Company shall maintain a centralized AI System Registry administered by the AI Governance Committee. No AI System may be used for Company business activities — including research, clinical, regulatory, legal, commercial, or administrative functions — unless it is listed in the Registry as Approved or Conditionally Approved for the intended use case.

4.2 Registry Entry Requirements

Each entry in the AI System Registry shall document, at minimum:

- Vendor name, AI System name, and Model Version(s) covered by the entry;
- Intended use case(s) and applicable department(s);
- Risk tier classification per Appendix A;
- GxP applicability determination (GxP-applicable / Non-GxP);
- PHI processing determination (PHI / No PHI);
- Data transmission analysis (data transmitted to external servers / on-premises / private cloud);
- Validation status and validation document references;
- Human oversight requirement(s) for the approved use case;
- Approval date and approving authority;
- Next review date;
- Executed Vendor Assessment (Section 13) reference.

4.3 New AI System Request Process

Any employee wishing to use an AI System not currently in the Registry shall submit an AI System Request Form (AISU-001) to the AI Governance Committee. The AGC shall evaluate the request within 30 calendar days and issue one of: Approved, Conditionally Approved (subject to specified controls), or Denied. Emergency requests from Regulatory Affairs in connection with time-sensitive submission activities may be evaluated on an expedited 10-business-day timeline with written justification.

Prohibited: Shadow AI

The use of AI Systems not listed in the Registry — commonly referred to as "shadow AI" — is strictly prohibited. This prohibition applies regardless of whether the AI System is accessed through a personal device, a personal account, or a free-tier subscription not associated with Company credentials. Employees who become aware of shadow AI use within the organization are required to report it to the Compliance Hotline.

Section 5 — Approved Use Cases by Department

The following table sets forth approved, conditionally approved, and prohibited AI use cases by department. "CONDITIONAL" designations require implementation of the specified controls before the AI System may be used. Use cases not listed must be submitted to the AGC for review before initiation.

Department	Use Case	Status	Oversight Requirement	GxP?
Research & Discovery	Scientific literature review, systematic review synthesis, and target validation report drafting	**PERMITTED**	PhD Scientist review and citation verification	No
Research & Discovery	Hypothesis generation and experimental design ideation (non-determinative)	**PERMITTED**	PI or senior scientist review	No
Research & Discovery	Analysis and summarization of internal research reports and data packages	**CONDITIONAL**	PhD Scientist review; CCI data handling controls apply	No
Research & Discovery	Drafting of laboratory notebook entries or recording of experimental results	**PROHIBITED**	N/A	**Yes**
Clinical Operations	Protocol template population and initial protocol draft generation	**CONDITIONAL**	Medical Director + Regulatory Affairs review; Validated Use required	**Yes**
Clinical Operations	Protocol amendment impact assessment and cross-reference identification	**CONDITIONAL**	Clinical Operations Lead + Regulatory review; GxP validation required	**Yes**
Clinical Operations	IRB/Ethics Committee submission cover letter drafting	**PERMITTED**	Clinical Operations Manager review	No
Clinical Operations	CRF annotation and data query response drafting from structured data	**CONDITIONAL**	Data Manager review; no PHI through external AI APIs	**Yes**
Regulatory Affairs	CTD Module 2 summary first-draft generation from structured study data	**CONDITIONAL**	Regulatory Writer + RA Director review; Validated Use; full attestation required	**Yes**
Regulatory Affairs	IND, NDA, BLA cross-reference and consistency checking	**CONDITIONAL**	Regulatory Writer review; Validated Use; all errors corrected before filing	**Yes**
Regulatory Affairs	FDA/EMA meeting request and correspondence drafting	**PERMITTED**	RA Director review and approval	No
Regulatory Affairs	Labeling precedent research and comparative label analysis	**PERMITTED**	Regulatory Affairs review	No
Legal / IP	Patent landscape monitoring and claim summary generation	**PERMITTED**	IP Analyst or Patent Attorney review; no FTO opinions from AI alone	No
Legal / IP	Contract review for risk flagging and clause summarization	**PERMITTED**	Attorney review required before execution	No
Legal / IP	Freedom-to-Operate analysis support (claim mapping only)	**CONDITIONAL**	Registered Patent Attorney renders all legal opinions; AI provides research support only	No
Pharmacovigilance	ICSR narrative drafting from structured CRF/safety database data	**CONDITIONAL**	PV Physician review and approval; no PHI via external API; Validated Use	**Yes**

Pharmacovigilance	DSUR aggregate safety section drafting and signal trend narrative	**CONDITIONAL**	PV Medical Director review; Validated Use required	**Yes**
Commercial / Med Affairs	Internal market landscape research summarization	**PERMITTED**	MSL or Medical Affairs review	No
Commercial / Med Affairs	Draft scientific exchange communications (internal use only)	**PERMITTED**	Medical Affairs physician review before any external use	No
All Departments	Meeting minutes and summary drafting from transcripts or notes	**PERMITTED**	Meeting organizer review and approval	No
All Departments	Training material and SOP drafting (non-GxP SOPs)	**PERMITTED**	Department head review and QMS approval per standard SOP workflow	No

Section 6 — Prohibited Uses

The following uses of AI Systems are prohibited without exception. These prohibitions apply regardless of whether the AI System has been approved for other use cases, regardless of time pressures or business justification, and regardless of whether the prohibited use produces an output that appears accurate. Violations are reportable events subject to Section 15 (Consequences).

6.1 Data Integrity and Regulatory Prohibitions

- **Fabrication or Alteration of Scientific Data:** Using AI to generate, alter, or extrapolate experimental data, analytical results, biomarker values, clinical outcome data, or any other scientific data that is represented as empirically derived. Example: Using an LLM to generate plausible-appearing pharmacokinetic parameters for a study that was not conducted, or to smooth statistical outliers without proper statistical justification.
- **Unreviewed AI Content in Regulatory Submissions:** Incorporating AI-generated text, tables, or analyses into any IND, NDA, BLA, IND amendment, NDA supplement, DSUR, PSUR, or other FDA or international regulatory agency submission without the review, editing, and attestation of a qualified regulatory professional. Example: Copying AI-generated Module 2.7.3 clinical summary text into a submission document without substantive expert review.
- **AI-Generated Legal Opinions or FTO Conclusions:** Representing AI-generated patent analysis, freedom-to-operate assessment, or legal conclusion as a legal opinion, or relying on such output as a substitute for a signed attorney opinion when making patent clearance or business development decisions.
- **AI Authorship of GLP Study Records:** Using AI to generate or populate raw data records, study plan entries, deviation forms, audit reports, or any other record subject to 21 CFR Part 58 (GLP). GLP records must reflect actual observations made by qualified study personnel and cannot be delegated to AI.
- **Clinical Data Fabrication or Imputation:** Using AI to generate, impute, or infer individual subject clinical data (e.g., adverse event severity grades, laboratory values, vital signs, efficacy assessments) that was not actually collected from the subject in accordance with the approved protocol.

6.2 Data Privacy Prohibitions

- **PHI Transmission to External AI APIs:** Transmitting Protected Health Information — including subject-level adverse event narratives, medical history summaries, genomic data, or any other individually identifiable health information — to any external AI system via API or browser interface, regardless of whether a Business Associate Agreement is in place, unless the AI system is deployed within a HIPAA-compliant infrastructure meeting the requirements of 45 CFR § 164.312 and has been approved for PHI processing by the Privacy Officer.
- **Transmission of Unrestricted CCI to Consumer AI Platforms:** Inputting unpublished compound structures, synthesis routes, clinical trial results, regulatory strategy documents, business development term sheets, or attorney-client privileged communications into consumer-tier AI platforms (e.g., free-tier ChatGPT, Gemini, or Claude.ai personal accounts) where the vendor's terms permit use of inputs for model training or do not provide adequate confidentiality protections.

6.3 Operational and Ethical Prohibitions

- **Undisclosed AI Authorship in Patient-Facing or Investigator-Facing Communications:** Representing AI-generated content as solely the work of a named individual in any communication directed to clinical trial investigators, patients, or IRB/Ethics Committees without disclosure that AI assisted in drafting.
- **Circumvention of Safety Monitoring Requirements:** Using AI to make or defer safety reporting decisions — including determinations of expectedness, seriousness, or causality for adverse events — without the mandatory pharmacovigilance physician review required by 21 CFR § 312.32 and ICH E2A.
- **Shadow AI Use:** Using any AI System not listed in the Company's AI System Registry for any Company business purpose, including research, drafting, analysis, or communication tasks. This prohibition applies to personal AI subscriptions and free-tier AI tools accessed on personal devices when used for Company business.

Section 7 — Data Classification and Handling Requirements

All data processed through AI Systems must be classified in accordance with the Company's Data Classification Framework. The following table governs AI-specific handling requirements for each data classification tier. These requirements supplement, and do not replace, the baseline data handling requirements applicable under the Company's Information Security Policy.

Data Classification	Risk Level	Examples	AI-Specific Handling Requirements
PHI / ePHI	**CRITICAL**	Subject-level clinical data, adverse event narratives with subject identifiers, signed ICFs, genomic data linked to subject IDs	Prohibited from external AI APIs without HIPAA-compliant infrastructure; on-premises or private cloud deployment required; BAA required; encryption at rest and in transit; access limited to need-to-know; PO approval required before any AI processing
Confidential Company Information (CCI) — Tier 1	**HIGH**	Unpublished compound structures, IND/NDA/BLA draft content, clinical trial efficacy/safety results, regulatory strategy, M&A targets	Enterprise AI platforms with executed DPA/BAA only; no consumer AI platforms; no training-data opt-in; data residency controls; reviewed and approved by IT Security and Legal before AI use
CCI — Tier 2	**MEDIUM**	Internal SOPs, business plans, vendor agreements, HR data, budget information	Enterprise platforms only; standard DPA protections; internal use only; no transmission to systems permitting training-data opt-in
Internal Use Only	**LOW**	Meeting agendas, internal project plans, non-sensitive correspondence, publicly available literature summaries	Approved AI platforms; standard access controls; no special restrictions beyond Registry requirement
Public Information	**NONE**	Published scientific literature, publicly filed patent documents, FDA guidance documents, press releases	No restrictions on AI processing beyond Registry requirement

7.1 PHI-Specific Requirements

The Company's use of PHI in AI workflows is subject to the HIPAA Privacy Rule (45 CFR Part 164, Subpart E) and Security Rule (45 CFR Part 164, Subpart C). The following controls are mandatory for any AI application that processes PHI:

- A fully executed Business Associate Agreement (BAA) with the AI System vendor is required prior to any PHI processing. The BAA must specify the permitted and required uses and disclosures of PHI, the vendor's obligation to implement appropriate safeguards, and the vendor's obligation to report unauthorized disclosures.
- Technical safeguards per 45 CFR § 164.312 must be implemented, including access controls (§ 164.312(a)(1)), audit controls (§ 164.312(b)), integrity controls (§ 164.312(c)(1)), and transmission security (§ 164.312(e)(1)).
- De-identification of PHI prior to AI processing is the preferred approach where the research or operational objective can be achieved without subject identifiers. De-identification must comply with either the Safe Harbor method (45 CFR § 164.514(b)(2)) or the Expert Determination method (45 CFR § 164.514(b)(1)).
- Any use of PHI in AI model fine-tuning — including fine-tuning conducted by or on behalf of the Company — requires prior IRB review of the de-identification protocol and written approval from the Privacy Officer.

Section 8 — GxP Validation Requirements

8.1 Validation Framework

AI Systems used in GxP Contexts are subject to the computerized system validation requirements applicable under 21 CFR Parts 11, 58, 210-211, and 312, and are governed by the ISPE GAMP 5 Second Edition (2022) framework for Good Automated Manufacturing Practice. Under GAMP 5, AI Systems incorporating configurable LLMs are classified as Category 5 software and require the full validation lifecycle described in this Section.

Validation of an AI System does not mean that the AI System produces uniformly correct outputs — LLMs are inherently probabilistic and non-deterministic. Validation demonstrates that: (a) the system performs its intended function consistently within defined parameters; (b) its known limitations are understood and controlled through human oversight and workflow design; and (c) the combination of the AI System and its human oversight controls is fit for the intended GxP purpose.

8.2 Validation Lifecycle Requirements

8.2.1 User Requirements Specification (URS)

A URS shall be prepared for each GxP-applicable AI use case, defining: the intended use and scope; functional requirements; performance requirements with acceptance criteria; security and access control requirements; data integrity requirements; interface requirements with upstream and downstream systems; and exclusions from scope. The URS is the foundation against which validation testing is designed.

8.2.2 Risk Assessment

A risk assessment shall be conducted for each GxP AI application prior to validation, assessing the potential for patient harm, data integrity compromise, or regulatory non-compliance resulting from system failure or incorrect output. The risk assessment shall determine the required rigor of validation testing (IQ only / IQ+OQ / full IQ+OQ+PQ) in accordance with the risk-tiering framework in Appendix A.

8.2.3 Installation Qualification (IQ)

IQ testing shall verify that the AI System is installed in the correct configuration, in the approved environment, with the specified software version (including Model Version), with appropriate access controls, and with audit trail functionality enabled. IQ documentation shall record the Model Version as a controlled parameter.

8.2.4 Operational Qualification (OQ)

OQ testing shall verify that the AI System performs its specified functions across a representative range of test cases, including expected use cases, boundary conditions, and adversarial inputs. For LLM-based systems, OQ shall include testing for: accuracy of factual claims against verified sources (citation verification testing); appropriate handling of out-of-scope queries; output consistency across repeated identical prompts; and appropriate escalation or refusal behavior for prohibited use case categories. Acceptance criteria shall be pre-specified in the OQ protocol.

8.2.5 Performance Qualification (PQ)

PQ testing, required for all CRITICAL-risk AI applications, shall demonstrate that the system performs consistently in the production environment over a defined monitoring period (minimum 90 days) using representative production data. PQ shall include statistical sampling of AI outputs for quality review against a defined acceptance rate threshold. PQ acceptance must be documented prior to production release.

8.2.6 Change Control

Any change to a validated AI System — including but not limited to a change in Model Version (including vendor-pushed updates), a change in system configuration, a change in integration with other systems, or a change in intended use —

shall be assessed through the Company's change control SOP (SOP-CHNG-001) prior to implementation. The change assessment shall determine whether re-qualification (IQ/OQ/PQ, in full or in part) is required.

Critical: Model Version Changes

Enterprise LLM providers including Anthropic and OpenAI periodically update underlying model weights, which constitutes a change in the validated system. The Company must contractually require advance notice (minimum 90 calendar days) of planned model version changes from AI System vendors. Internal IT must monitor for unannounced changes. Any model version change for a validated AI System triggers mandatory change control assessment before continued GxP use.

8.3 Ongoing Performance Monitoring

Following PQ acceptance, validated AI Systems shall be subject to ongoing performance monitoring including: (a) periodic sampling review of AI outputs incorporated into GxP documents, with deficiency rates tracked against pre-specified thresholds; (b) quarterly review by the AI Governance Committee of operational metrics, complaint reports, and deviation records; and (c) annual re-validation assessment to determine whether re-qualification is warranted by accumulated changes or performance data.

Section 9 — Human Oversight Requirements

9.1 Principle of Meaningful Oversight

The Company requires that human oversight of AI outputs be substantive, not perfunctory. A reviewer who signs an AI-assisted document has exercised professional judgment about the accuracy, completeness, and appropriateness of the content — including AI-generated content — and bears professional responsibility for that judgment. Reviewers are expected to identify and correct errors, hallucinations, unsupported claims, and regulatory non-compliance in AI outputs before those outputs are incorporated into any document. Reviewers who sign AI-assisted documents without substantive review are in violation of this Policy and, in GxP contexts, may be in violation of applicable federal regulations.

9.2 Oversight Requirements by Use Case

Use Case	GxP?	Required Reviewer Qualification	Minimum Review Obligations	Risk Tier
Literature review synthesis / target validation report	No	PhD Scientist in relevant discipline	Citation verification (100%); scientific accuracy review; sign-off in DMS	**LOW**
Protocol first draft generation	**Yes**	MD or equivalent clinical lead + Regulatory Affairs professional	Line-by-line review; consistency check against IND; eTMF documentation	**HIGH**
Protocol amendment impact assessment	**Yes**	Clinical Operations Lead + Regulatory Affairs	Cross-reference verification; QA sign-off; change control record	**HIGH**
CTD Module 2 summary drafting	**Yes**	Regulatory Affairs professional (RAC or equivalent) + RA Director	Line-by-line review; all data citations independently verified against source studies; attestation memo filed	**CRITICAL**
ICSR narrative drafting	**Yes**	Pharmacovigilance Physician (MD with PV experience)	Medical accuracy review; MedDRA coding verification; causality re-assessment; physician signature in safety database	**CRITICAL**
DSUR safety signal narrative	**Yes**	PV Medical Director	Signal assessment review; comparison against cumulative safety database; regulatory strategy alignment	**CRITICAL**
Patent claim map / FTO support	No	Registered Patent Attorney (USPTO)	Legal analysis conducted independently; AI output used as research starting point only; attorney signs FTO opinion	**MEDIUM**
Contract review / risk flagging	No	Attorney (Company Legal or outside counsel)	Legal review before execution; AI flags reviewed against full contract text	**MEDIUM**
Meeting minutes drafting	No	Meeting organizer or designated minute-taker	Accuracy review against recording/notes; approval before distribution	**LOW**
Non-GxP SOP drafting	No	Department head + QA (per SOP approval workflow)	Standard QMS review and approval process; no additional AI-specific requirement	**LOW**

9.3 Citation Verification Requirement

AI Systems — including state-of-the-art LLMs — are known to produce hallucinated citations: references that appear facially plausible but that either do not exist or do not support the attributed claim. Accordingly, any AI-generated document that incorporates citations to scientific literature, regulatory precedent, patent records, or other external sources shall be subject to 100% citation verification by the reviewing professional prior to use. Verification means confirming that: (a) the cited source exists and is retrievable; and (b) the cited source substantiates the claim attributed to it. Use of AI-generated content with unverified citations in regulated documents is a prohibited practice.

Section 10 — Audit Trail and Documentation Requirements

10.1 Regulatory Basis

21 CFR § 11.10(e) requires that computerized systems used in the creation, modification, maintenance, archival, retrieval, or transmission of electronic records that are subject to FDA regulations shall use computer-generated, time-stamped audit trails to independently record the date and time of operator entries and actions that create, modify, or delete electronic records. The use of AI Systems to generate content that is incorporated into regulated electronic records extends these requirements to capture the provenance and human review of AI-generated content.

10.2 Mandatory Audit Trail Elements

For any AI-generated content incorporated into a GxP-applicable document, the Audit Trail shall capture the following elements at a minimum. These elements shall be stored in the Company's validated document management system (DMS) or equivalent validated electronic repository, and shall not be modifiable after creation without generation of a new audit trail entry:

Audit Trail Element	Description and Requirements
AI Session Identifier	A unique identifier assigned to each AI interaction session, enabling reconstruction of the complete chain of prompts and responses associated with a given output.
AI System Name and Version	The name of the AI System (e.g., Claude 3.5 Sonnet, GPT-4o) and the Model Version identifier as provided by the vendor. If the vendor does not provide a version identifier accessible to the user, this shall be documented and escalated to IT for resolution.
Generation Timestamp	UTC timestamp of AI output generation, accurate to the second.
Operator Identity	The authenticated identity of the Company employee or contractor who submitted the prompt and received the output. Shared accounts are prohibited for any GxP AI interaction.
Prompt Reference	Where the Prompt contains no CCI or PHI, the full Prompt text shall be captured. Where the Prompt contains CCI or PHI, a hash of the Prompt together with a reference to the stored Prompt (maintained with appropriate access controls) shall be captured.
Output Reference	A reference to the stored AI output (full text or document location). Where the output is directly incorporated into a document, the document identifier and version shall be recorded.
Reviewer Identity	The authenticated identity of the qualified professional who reviewed and approved the AI-generated content for use in the regulated document.
Review Timestamp	UTC timestamp of human review and approval, accurate to the second.
Material Changes Description	A description (not necessarily exhaustive) of material changes made by the reviewer to the AI output before incorporation — e.g., "Corrected two hallucinated citations; revised regulatory precedent analysis; adjusted dosing language for consistency with IND."
Document Reference	The identifier of the final regulated document into which the AI-generated content was incorporated, including document version.

10.3 Retention Requirements

Audit trail records for AI use in GxP Contexts shall be retained for the longer of: (a) the retention period applicable to the underlying regulated document under the applicable FDA regulation (e.g., 15 years for clinical development records under 21 CFR § 312.62(b)); or (b) 10 years from the date of generation. Audit trail records shall be maintained in a format that enables their accurate and ready retrieval throughout the retention period, per 21 CFR § 11.10(c).

10.4 Non-GxP Documentation

For AI use in non-GxP contexts, the documentation requirements are lighter but not absent. At minimum, employees shall retain records sufficient to identify: the AI System used; the general nature of the task for which AI was used; and the name of the human reviewer who approved the final output. These records may be maintained in the Company's standard document management system or project management platform.

Section 11 — Intellectual Property and Copyright Considerations

11.1 Copyright Status of AI-Generated Output

The U.S. Copyright Office has issued guidance — most recently in its March 2023 guidance following Thaler v. Vidal and the Zarya of the Dawn registration decision, and its subsequent February 2023 policy statement — establishing that purely AI-generated content (i.e., content generated without sufficient human creative authorship) is not eligible for copyright protection under 17 U.S.C. § 102. The copyright interest in AI-assisted works derives from the human creative contribution to the final output, including the creative exercise of selection, arrangement, and expression in prompting, reviewing, and editing AI-generated content.

Practical implications for the Company: (a) AI-generated regulatory submission text, as substantially edited and approved by qualified regulatory professionals, is likely to carry Company copyright in the human-authored elements, but the extent of protection is fact-specific; (b) minimally edited AI output incorporated directly into Company documents may not receive copyright protection, creating freedom-to-operate risks if the text is reproduced by third parties; and (c) the Company shall ensure that all significant AI-assisted documents reflect substantive human authorship to maximize copyright protection.

11.2 Training Data and Third-Party Copyright

The Company acknowledges ongoing legal uncertainty regarding: (a) whether the training of LLMs on copyrighted text constitutes copyright infringement (see Tremblay v. OpenAI, Inc., N.D. Cal. 2023; Andersen v. Stability AI Ltd., N.D. Cal. 2023); and (b) whether AI outputs that closely replicate training data may themselves infringe third-party copyrights. The Company takes the following positions to manage this risk:

- Employees shall not direct AI Systems to reproduce substantial portions of copyrighted works (including journal articles, textbooks, patent claims, or proprietary databases) verbatim, and shall report any AI output that appears to reproduce verbatim text from identifiable copyrighted sources to the Legal Department.
- All AI-generated content incorporated into regulatory submissions, scientific publications, patent applications, or commercial materials shall be reviewed by Legal or outside counsel for copyright clearance risk before filing or public disclosure.
- The Company shall maintain records of the AI Systems used for each significant work product to enable assessment of training data provenance to the extent vendor documentation permits.

11.3 Patent Implications of AI-Assisted Invention

The United States Patent and Trademark Office (USPTO) has confirmed through its February 2024 guidance on AI-assisted inventions that AI systems cannot be listed as inventors on U.S. patent applications; only natural persons who make significant contributions to the conception of the claimed invention qualify as inventors under 35 U.S.C. § 100(f). AI-assisted inventions where human inventors make significant contributions to the conception of the invention remain patentable.

Employees are required to disclose to the Legal/IP Department any invention or potential invention where an AI System made a material contribution to the conception of the invention, so that: (a) inventorship can be properly analyzed; (b) the appropriate inventor oath/declaration can be prepared; and (c) the invention record accurately reflects the role of AI in the inventive process. Failure to disclose AI involvement in inventive activities may give rise to inequitable conduct issues before the USPTO.

11.4 Vendor AI Output License Terms

Employees shall not use AI Systems in connection with Company business activities unless the applicable vendor terms of service or enterprise agreement: (a) confirm that the Company owns or has a license to AI-generated outputs sufficient for the intended business purpose; (b) do not require the Company to grant the vendor any rights in Company inputs or outputs that are inconsistent with the Company's confidentiality and IP obligations; and (c) do not permit the vendor to use Company inputs as training data without explicit Company consent. Legal shall review vendor AI terms before execution of any enterprise AI agreement.

Section 12 — Vendor Assessment Requirements

No AI System vendor may be added to the AI System Registry until the AI Governance Committee has completed a Vendor Assessment evaluating the criteria set forth below. Vendor Assessments shall be conducted by IT Security in collaboration with Legal and, for GxP-applicable systems, Quality Assurance. Assessments shall be repeated at each enterprise contract renewal and upon any material change in the vendor's ownership, security certifications, or data handling practices.

Assessment Criterion	Requirement Detail	Classification
Security Certification	SOC 2 Type II report (current, within 12 months); ISO 27001 certification preferred for vendors processing CCI or PHI	**REQUIRED**
HIPAA BAA Availability	Vendor must be willing and able to execute a BAA meeting 45 CFR § 164.504(e) requirements for any PHI-processing use case	**REQUIRED**
Data Processing Agreement	GDPR-compliant DPA for any data involving EU subjects; CCPA-compliant terms for California residents	**REQUIRED**
Training Data Opt-Out	Vendor must commit in writing that Company inputs will not be used to train or fine-tune the underlying model without explicit Company consent	**REQUIRED**
Model Version Control	Vendor must provide advance notice (minimum 90 days) of planned model version changes for enterprise customers	**REQUIRED**
Data Residency	Ability to specify data residency (US-only or EU-only) for all Company data processed through the platform	**RECOMMENDED**
Penetration Testing	Annual third-party penetration testing results available to enterprise customers upon request	**RECOMMENDED**
Audit Rights	Contract includes right for Company (or designated auditor) to audit vendor's AI-related data handling practices	**RECOMMENDED**
FDA 21 CFR Part 11 Support	Vendor provides documentation supporting customer's Part 11 validation obligations (system description, IQ documentation support)	**RECOMMENDED**
Model Card / Transparency Documentation	Vendor publishes model cards or equivalent documentation describing model capabilities, limitations, known failure modes, and training data sources	**RECOMMENDED**
Incident Response SLA	Vendor maintains documented breach notification obligations meeting HIPAA's 60-day notification requirement (45 CFR § 164.410)	**RECOMMENDED**

Section 13 — Employee Training Requirements

13.1 Mandatory Training — All Employees

All employees and contractors subject to this Policy shall complete AI Governance Fundamentals training (Course AI-101) within 30 days of the Policy's effective date (for existing employees) or within 30 days of onboarding (for new hires and contractors). AI-101 covers: the scope and requirements of this Policy; the AI System Registry and approval process; data classification requirements for AI inputs; the prohibition on shadow AI; and how to report AI-related concerns. Completion of AI-101 is a condition of continued employment and contract engagement.

13.2 Role-Specific Training

Training Course	Required For	Key Content Areas
AI-101: Governance Fundamentals	All employees and contractors	Policy scope; Registry; data classification; prohibited uses; reporting obligations; shadow AI prohibition
AI-201: GxP AI Use and Validation	Research, Clinical Ops, Regulatory, QA, PV	21 CFR Part 11 requirements; GAMP 5 framework; validation lifecycle; change control; audit trail requirements; human oversight obligations in GxP contexts
AI-301: Regulatory Submission AI Use	Regulatory Affairs personnel	FDA AI/ML guidance (FDA-2023-N-3379; 2025 AI Action Plan); CTD AI-assisted drafting controls; attestation requirements; disclosure obligations; CRL risk management
AI-302: Pharmacovigilance AI Use	PV and Drug Safety personnel	ICH E2A/E2B/E2E requirements; ICSR narrative AI assistance controls; causality assessment responsibilities; MedDRA coding verification; DSUR AI-assisted drafting controls
AI-401: HIPAA and AI Privacy Compliance	Clinical Operations, PV, Biostatistics, Data Management	PHI definition and identification; de-identification methods; BAA requirements; prohibited PHI transmissions; breach reporting obligations; technical safeguards for AI-PHI systems
AI-501: IP and AI Governance	Legal, IP, Regulatory, Senior Scientists	Copyright status of AI outputs; inventorship analysis for AI-assisted inventions; USPTO guidance on AI-assisted inventions; training data copyright risk; vendor IP term assessment
AI-601: AI Governance Committee Training	AGC members	AGC charter and authority; use case evaluation framework; risk-tiering methodology; vendor assessment process; exception request evaluation; annual review process

13.3 Training Records and Competency Verification

All training completion records shall be maintained in the Company's Learning Management System (LMS) for a minimum of 5 years. Annual refresher training (minimum 1 hour) covering policy updates, new AI capabilities and risks, and lessons learned from internal AI quality events shall be required for all employees. Managers are responsible for ensuring that their direct reports complete required training within the specified timelines and shall escalate non-completion to HR and Compliance within 5 business days of the deadline.

Section 14 — AI Governance Committee

14.1 Composition

The AI Governance Committee (AGC) shall be constituted as a standing cross-functional committee with the following composition:

- Chair: Chief Compliance Officer (or designee)
- Members: VP Regulatory Affairs; General Counsel or Associate General Counsel; Chief Information Officer or VP of Informatics; Head of Quality Assurance; Chief Scientific Officer or VP Research; Head of Pharmacovigilance; Head of Clinical Operations
- Non-voting advisors (as needed): External regulatory counsel; IT Security; outside AI/ML technical advisors

14.2 Responsibilities

- Maintain and update the AI System Registry;
- Evaluate and approve or deny new AI System requests (AISU-001);
- Conduct annual Policy review and propose updates to the Board;
- Review and disposition AI-related deviations, quality events, and Compliance Hotline reports;
- Evaluate and grant or deny Policy exception requests;
- Monitor the regulatory environment (FDA, EMA, ICH, EU AI Act) for developments requiring Policy updates;
- Provide quarterly AI program status reports to the Board Audit Committee.

14.3 Meeting Cadence

The AGC shall meet at minimum monthly. Emergency meetings may be convened within 48 hours upon written request of the Chair, any two Members, or the CEO. AGC deliberations and decisions shall be documented in meeting minutes maintained in the Company's QMS.

Section 15 — Policy Violations and Consequences

Violations of this Policy are subject to disciplinary action calibrated to the severity of the violation, the degree of intent or recklessness involved, the harm or potential harm caused, and the individual's prior compliance history. The consequences set forth below are illustrative; actual consequences will be determined by HR, Legal, and the Chief Compliance Officer following investigation. Nothing in this Section limits the Company's right to take any additional action required by law, regulation, or contractual obligation.

Violation Type	Severity	Consequences
Shadow AI use for non-sensitive business tasks (single incident)	**LOW**	Documented counseling; mandatory repeat of AI-101 training within 14 days; notation in personnel record
Repeated shadow AI use or shadow AI use involving CCI-Tier 1 or PHI	**HIGH**	Formal written warning; mandatory AI-201 training; 90-day enhanced monitoring; escalation to HR and CCO
Transmission of PHI to non-approved AI system (inadvertent)	**HIGH**	Immediate incident report to Privacy Officer; HIPAA breach risk assessment; regulatory notification if required under 45 CFR § 164.410; formal written warning
Transmission of PHI to non-approved AI system (intentional or reckless)	**CRITICAL**	Immediate suspension pending investigation; potential termination; referral to HHS OCR if warranted; potential civil liability
Incorporation of unreviewed AI content into regulatory submission	**CRITICAL**	Immediate regulatory impact assessment by Regulatory Affairs and Legal; submission withdrawal or amendment if required; mandatory root cause analysis; potential termination; FDA notification if data integrity is implicated
Fabrication or alteration of scientific data using AI	**CRITICAL**	Immediate termination; referral to FDA (21 CFR § 312.70 / § 314.81(b)(1)(iii)); potential criminal referral to DOJ; civil litigation exposure; retraction of affected publications or submissions
Failure to disclose AI-assisted invention to Legal/IP	**MEDIUM**	Documented counseling; mandatory AI-501 training; IP audit of relevant invention records
Failure to complete mandatory AI training by deadline	**LOW**	Escalation to HR; training completion required within 7 business days; persistent non-completion may result in access restrictions to AI Systems

Reporting AI Policy Concerns

Employees who believe they have witnessed a violation of this Policy, or who have questions about whether a proposed AI use is compliant, are strongly encouraged to contact the Compliance Hotline (anonymous reporting available) or their direct manager, department head, or the Chief Compliance Officer. Retaliation against any employee who reports a good-faith compliance concern is strictly prohibited and is itself a terminable offense.

Section 16 — Annual Review Process

16.1 Review Trigger and Timeline

This Policy shall be formally reviewed no less than once every twelve (12) months from the Effective Date, or earlier upon any of the following triggers: (a) issuance of new or materially revised FDA guidance on AI in drug development; (b) new ICH guidelines addressing AI in GxP contexts; (c) material developments in EU AI Act implementing acts or associated technical standards; (d) a significant AI-related quality event, data breach, or regulatory finding at the Company; or (e) a material change in the Company's AI System portfolio or business operations.

16.2 Review Process

6. The AI Governance Committee shall initiate the review no later than 60 days before the anniversary of the Policy's effective date.
7. The AGC shall commission a structured review encompassing: (a) assessment of regulatory developments (FDA, EMA, ICH, EU AI Act) since the prior review; (b) analysis of internal AI quality events, training completion metrics, and Compliance Hotline reports; (c) benchmarking against industry peer AI governance practices; and (d) review of AI System Registry to confirm all listed systems remain current and appropriately characterized.
8. Proposed revisions shall be circulated to all AGC members and to key functional stakeholders (Department Heads of Research, Clinical, Regulatory, Legal, and QA) for comment. A 15-business-day comment period shall be provided.
9. The final revised Policy shall be presented to the Board Audit Committee for approval. Minor administrative changes (e.g., updating cross-references, correcting typographical errors, adding new AI Systems to the Registry) may be approved by the AGC Chair and CCO without full Board review.
10. The approved revised Policy shall be published to all employees within 5 business days of Board approval. Employees shall be notified of material changes and, where material changes affect training obligations, shall complete updated training within 45 days of notification.
11. A complete record of each annual review — including the review report, stakeholder comments, proposed and final revisions, and Board approval documentation — shall be retained in the Company's QMS for a minimum of 10 years.

Appendix A — AI System Risk-Tiering Framework

The following framework governs the assignment of risk tiers to AI System use cases. Risk tier assignment drives validation requirements, oversight requirements, and AGC review frequency. Tier assignment is made by the AI Governance Committee upon evaluation of an AISU-001 request and may be updated as use cases evolve.

Risk Tier	Tier Criteria	Governance Requirements	Example Use Cases
CRITICAL	GxP-applicable; PHI or CCI-Tier 1 processed; output directly incorporated into regulatory submission or safety report	Full IQ+OQ+PQ; ongoing PQ monitoring; AGC + QA approval; annual re-validation assessment; mandatory validated use designation; CRITICAL oversight requirements	Regulatory submission drafting; ICSR narrative generation; DSUR aggregate analysis; clinical protocol generation
HIGH	GxP-applicable; CCI processed; output incorporated into controlled GxP documents (protocols, IBs, CSRs)	IQ+OQ required; PQ recommended; AGC approval; quarterly performance review; HIGH oversight requirements	Protocol amendment management; IB drafting support; CSR section drafting
MEDIUM	Non-GxP but CCI-Tier 1 processed; or GxP-applicable with minimal regulatory impact; legal/IP functions	IQ+OQ required; AGC approval; semi-annual review; MEDIUM oversight requirements	Patent landscape analysis; FTO claim mapping; contract review; internal regulatory strategy drafting
LOW	Non-GxP; CCI-Tier 2 or Internal-Use-Only data; administrative or research-support functions	URS and risk assessment required; AGC approval; annual review; standard oversight	Literature review synthesis; meeting minutes; non-GxP SOP drafting; internal correspondence
MINIMAL	Non-GxP; public information only; no CCI, no PHI	Registry listing required; no formal validation; standard training completion	Public patent database searches; published literature summarization; public FDA guidance synthesis

Appendix B — Regulatory Reference Index

The following regulatory and guidance documents are referenced in this Policy. The most current version of each document shall govern. The AI Governance Committee shall monitor for updates and initiate Policy review as needed.

Document	Relevance to This Policy
21 CFR Part 11 (Electronic Records; Electronic Signatures)	Audit trail requirements; electronic record integrity; system validation for regulated records
21 CFR Part 58 (Good Laboratory Practice)	GLP requirements for AI use in nonclinical study records
21 CFR Parts 312, 314 (IND, NDA)	Clinical development and submission requirements; human oversight standard for sponsor responsibilities
21 CFR Part 601 (BLA)	Biologics licensing requirements; sponsor accountability for submission accuracy
ICH E6(R3) — Good Clinical Practice	GCP requirements applicable to clinical protocol development and clinical data management AI use
ICH E2A — Clinical Safety Data Management	Adverse event reporting requirements governing ICSR narrative AI assistance
ICH E2B(R3) — Electronic Transmission of Individual Safety Reports	ICSR format and content standards applicable to AI-assisted narrative drafting
ISPE GAMP 5 Second Edition (2022)	Computerized system validation framework; Category 5 classification

	for AI/ML systems
HIPAA Privacy Rule — 45 CFR Part 164, Subparts A & E	PHI definition; minimum necessary standard; de-identification methods; BAA requirements
HIPAA Security Rule — 45 CFR Part 164, Subpart C	Technical safeguards for ePHI processed through AI systems
EU AI Act — Regulation (EU) 2024/1689	Risk classification framework for AI systems; high-risk AI obligations; transparency requirements
FDA Discussion Paper — AI in Drug Development (Oct 2023, Docket FDA-2023-N-3379)	FDA's current policy position on sponsor use of AI in regulated activities
FDA AI Action Plan (January 2025)	FDA's regulatory agenda for AI governance in drug and device development through 2026
USPTO Guidance on AI-Assisted Inventions (February 2024)	Inventorship requirements; natural person requirement for patent inventors
U.S. Copyright Office AI Policy (2023)	Copyright status of AI-generated content; human authorship requirement
35 U.S.C. § 100(f) — Definition of Inventor	Patent law basis for natural-person-only inventorship

Appendix C — AI System Request Form (AISU-001) — Template

Form AISU-001

Submit to: aig-committee@[company].com | Processing time: 30 calendar days standard / 10 business days expedited

Field	Response (Complete All Fields)
Requestor Name and Title	
Department	
Date of Request	
AI System Name and Vendor	
Model Version (if known)	
Proposed Use Case Description	
Proposed Use Department(s)	
Is this use GxP-applicable? (Yes/No — explain)	
Will PHI be processed? (Yes/No — explain)	
Will CCI be processed? (Yes/No — classification tier)	
Will data be transmitted to external servers? (Yes/No)	
Proposed human oversight reviewer and qualifications	
Business justification and urgency	

Has a Vendor Assessment been completed? (Yes/No/Pending)	
Requested risk tier (with justification)	
Requestor Signature and Date	
Manager Approval Signature and Date	

Appendix D — Human Oversight Attestation Template

The following attestation shall be completed by the reviewing professional and filed in the document management system as a record associated with any regulated document incorporating AI-generated content.

Attestation Field	Entry
Document Title and Identifier	
Document Version	
AI System(s) Used	
Model Version(s)	
Date(s) of AI Generation	
AI Session Identifier(s)	
Sections of Document Containing AI-Generated Content	
Description of Review Conducted	
Material Corrections Made to AI Output	
Citations Verified (Y/N; attach verification log if applicable)	
Reviewer Name and Credentials	
Reviewer Signature	
Date of Review and Attestation	
Second Reviewer (if required by risk tier)	
Second Reviewer Signature and Date	

ATTESTATION STATEMENT: I attest that I have reviewed the AI-generated content identified above; that I possess the professional qualifications and subject matter expertise required to evaluate this content in its intended regulatory, scientific, or legal context; that I have made all corrections necessary to ensure the accuracy, completeness, and regulatory compliance of the final document; and that I take full professional and personal responsibility for the accuracy of the content of this document, including AI-assisted portions, as submitted or filed.

Appendix E — Glossary of Abbreviations

Abbreviation	Full Form
AGC	AI Governance Committee
AISU	AI System Use Request (Form AISU-001)
BAA	Business Associate Agreement (HIPAA)
BLA	Biologics License Application
CCI	Confidential Company Information
CCO	Chief Compliance Officer
CRF	Case Report Form
CRL	Complete Response Letter
CSR	Clinical Study Report
CTD	Common Technical Dossier (ICH M4)
DMS	Document Management System
DPA	Data Processing Agreement (GDPR)
DSUR	Development Safety Update Report (ICH E2F)
eTMF	Electronic Trial Master File
FDA	U.S. Food and Drug Administration
FTO	Freedom to Operate
GAMP	Good Automated Manufacturing Practice (ISPE)
GCP	Good Clinical Practice
GLP	Good Laboratory Practice
GMP	Good Manufacturing Practice
GxP	Good [any] Practice (collective GLP/GCP/GMP/GVP)
HIPAA	Health Insurance Portability and Accountability Act
IB	Investigator's Brochure
ICH	International Council for Harmonisation
ICSR	Individual Case Safety Report
IND	Investigational New Drug Application
IP	Intellectual Property
IQ	Installation Qualification
IRB	Institutional Review Board
LLM	Large Language Model
LMS	Learning Management System
MedDRA	Medical Dictionary for Regulatory Activities
NDA	New Drug Application
OQ	Operational Qualification

PHI	Protected Health Information
PO	Privacy Officer
PQ	Performance Qualification
PV	Pharmacovigilance
QA	Quality Assurance
QMS	Quality Management System
RAC	Regulatory Affairs Certification
URS	User Requirements Specification
USPTO	United States Patent and Trademark Office